

# Homework 2

José Javier González Ortiz

September 29, 2015

## 1 Efficiency Analysis

Given the Serial program we can analyze its Big-O notation. From the given information about initialize, statement 3, statement 5 and finalize we can calculate the serial time as follows:

$$\text{SerTime}(n) = \Theta(n) + n \cdot (\Theta(1) + n\Theta(1)) + \Theta(1) = \Theta(n^2) \quad (1.1)$$

Since we can perfectly parallelize the inner loop the Parallel time without any communication will be:

$$\text{ParTime}(n, p) = \Theta(n) + n \cdot \left( \Theta(1) + \frac{n}{p} \Theta(1) \right) + \Theta(1) = \Theta\left(\frac{n^2}{p}\right) + \Theta(n) \quad (1.2)$$

We cannot discard the  $\Theta(n)$  term since for  $p = \Theta(n)$  will as important as the first term. Introducing communication into the algorithm yields the following efficiencies:

- a) Introducing a  $\Theta(1)$  communication in the inner loop produces no changes in the big O representation of the Parallel Time

$$\text{ParTime}(n, p) = \Theta(n) + n \cdot \left( \Theta(1) + \frac{n}{p} \Theta(1) \right) + \Theta(1) = \Theta\left(\frac{n^2}{p}\right) + \Theta(n) \quad (1.3)$$

The efficiency will be:

$$\text{Efficiency}(n, p) = \frac{\Theta(n^2)}{p \left( \Theta\left(\frac{n^2}{p}\right) + \Theta(n) \right)} = \frac{\Theta(n^2)}{\Theta(n^2) + \Theta(np)} \quad (1.4)$$

This will only be constant as long as  $np = O(n^2)$ , namely if  $p = O(n)$ . If  $p = \omega(n)$  the efficiency will drop to zero, however we know by definition that  $p \leq n$  since there cannot be more processors than iterations of the inner loop.

- b) Introducing a  $\Theta(p)$  time in the outer loop the Parallel time will be:

$$\text{ParTime}(n, p) = \Theta(n) + n \cdot \left( \Theta(1) + \frac{n}{p} \Theta(1) + \Theta(p) \right) + \Theta(1) = \Theta\left(\frac{n^2}{p}\right) + \Theta(np) \quad (1.5)$$

With this expression the efficiency is:

$$\text{Efficiency}(n, p) = \frac{\Theta(n^2)}{p \left( \Theta\left(\frac{n^2}{p}\right) + \Theta(np) \right)} = \frac{\Theta(n^2)}{\Theta(n^2) + \Theta(np^2)} \quad (1.6)$$

As before, this ratio will tend to a constant given that  $np^2 = O(n^2)$ . Simplifying the expression we get  $p = O(\sqrt{n})$ . For  $p = \omega(\sqrt{n})$  the efficiency will go to zero.

c) Introducing a  $\Theta(\sqrt{n/p})$  time in the outer loop the Parallel time will be:

$$\text{ParTime}(n, p) = \Theta(n) + n \cdot \left( \Theta(1) + \frac{n}{p} \Theta(1) + \Theta\left(\sqrt{\frac{n}{p}}\right) \right) + \Theta(1) = \Theta\left(\frac{n^2}{p}\right) + \Theta(n) + \Theta\left(n\sqrt{\frac{n}{p}}\right) \quad (1.7)$$

The efficiency in this case follows a somewhat more complicate expression

$$\text{Efficiency}(n, p) = \frac{\Theta(n^2)}{p \left( \Theta\left(\frac{n^2}{p}\right) + \Theta(n) + \Theta\left(n\sqrt{\frac{n}{p}}\right) \right)} = \frac{\Theta(n^2)}{\Theta(n^2) + \Theta(np) + \Theta(\sqrt{n^3 p})} \quad (1.8)$$

To get a better idea of the asymptotic behavior we can divide by  $\Theta(n^2)$  and get

$$\text{Efficiency}(n, p) = \frac{\Theta(1)}{\Theta(1) + \Theta\left(\frac{p}{n}\right) + \Theta\left(\sqrt{\frac{p}{n}}\right)} \quad (1.9)$$

From this we have to satisfy both:

$$\frac{p}{n} = \Theta(1) \quad \sqrt{\frac{p}{n}} = \Theta(1) \quad (1.10)$$

Which will only happen if  $p = O(n)$ . Then the efficiency will tend to a constant. As we said before  $p \leq n$  so this will always be satisfied

## 2 Reduce Sum

a) We want to minimize the communication in order to improve the efficiency, so the best approach will be summing the  $n/p$  values in each processor and then adding these intermediate sums. Since they are arranged in a 1-dimensional mesh the running time of the communication is  $\Theta(p)$ . The procedure would be something similar to the one described in Algorithm 1.

---

### Algorithm 1 1D reduce(sum,A)

---

**Require:** Each processor has the common values  $p$  and  $n$ , as well as the local value  $id$

---

```

1: procedure REDUCE(SUM,A)
2:    $s \leftarrow \text{sum}[A(id \cdot n/p : (id + 1) \cdot n/p)]$ 
3:   if  $id > 0$  then
4:      $\text{receive}(id - 1, t)$ 
5:      $s \leftarrow s + t$ 
6:   if  $id < p - 1$  then
7:      $\text{send}(id + 1, s)$ 
8:   return  $s$ 
```

---

b) As we can see from the previous question the algorithm will take  $\Theta(n/p) + \Theta(p)$  where  $1 \leq p \leq n$ . It is easy to see that for both  $p = \Theta(1)$  and  $p = \Theta(n)$  the running time is  $\Theta(n)$ . We want to minimize so by making the terms being equal we will minimize the running time.

$$\Theta\left(\frac{n}{p}\right) = \Theta(p) \rightarrow p = \Theta(\sqrt{n}) \quad (2.1)$$

Therefore the best running time will happen when  $p = \Theta(\sqrt{n})$  and will be  $\text{ParTime} = \Theta(\sqrt{n})$

- c) In a fully connected computer, given a binary operator the best communication will be  $\Theta(\log p)$ . In this case the running time is going to be  $\Theta(n/p) + \Theta(\log p)$ .

Since due to the communication constraint we know that there is a  $\Omega(\log n)$  lower bound we can see that we cannot do better than getting rid of the linear term. This happens for  $p = \Theta(n)$  which minimizes the time to  $\text{ParTime} = \Theta(\log n)$

### 3 SIMD Hypercube

#### 3.1 Analysis

As a first step is good to think about which is the minimum ideal runtime given the communication constraints of an hypercube. Let  $p_M$  be the processor with the maximum value  $v_M \in V : v_M \geq v_i \quad \forall v_i \in V$ . There will exist a processor  $\overline{p_M}$  such that the minimum distance from  $p_M$  measured in edges will be  $\log p$ . To reason this, just think about this maximum point as a vertex in the hypercube and the processor  $\overline{p_M}$  as the opposite vertex, since the hypercube has  $\log p$  dimensions there is no shortest path between them. Furthermore, the coordinates  $\overline{p_M}$  will be the bitwise complement of the coordinates of  $p_M$ .

Thus the algorithm will only be correct if  $v_{\overline{p_M}} = v_M$  which will need at least  $\Omega(\log p)$  communications. We have found a logarithmic lower bound in our running time. An efficient algorithm that solves the problem can be thought in the following manner:

In each timestep  $i$ , each processor exchanges its value with its neighbor in the  $i^{\text{th}}$  dimension and stores as its new value the maximum of the two. After  $\log p$  steps all the processors have the maximum value in the array.

This algorithm is described in the Algorithm 2.

#### 3.2 Algorithm

---

##### Algorithm 2 SIMD Hypercube AllReduce(max)

---

**Require:** Each processor has the common value  $p$  and the local values  $pid$  and  $V$ , as well as the variable  $M$  and  $R$

The controller has the variables  $dim$  and  $x$

1: <b>procedure</b> ALLREDUCE(max, $V(0 : p - 1)$ )	
2: $M \leftarrow V$	This instruction is transmitted and executed in the processors
3: $x \leftarrow \lg p$	This instruction is executed in the Controller
4: <b>for</b> $dim = 0$ <b>to</b> $x - 1$ <b>do</b>	This instruction is executed in the Controller
5: $exchange(dim, M, R)$	This instruction is transmitted and executed in the processors
6: <b>if</b> $R > M$ <b>then</b>	This instruction is transmitted and executed in the processors
7: $M \leftarrow R$	This instruction is transmitted and executed in the processors
8: <b>return</b> $M$	This instruction is transmitted and executed in the processors

---

### 3.3 Correctness

For correctness first we state that given any processor  $p$  it has only one neighbor in the dimension  $dim$ , namely  $p \oplus 2^{dim}$ . At each time step the number of different values in the array is at least halved because the processors are saving the maximum value and each processor is *exchanging* the value with only one other processor. After  $\log p$  steps we will have gone through all the dimensions and all the processors will have the maximum value  $v_M$ .

Namely, let the  $p_M$  be the processor with the maximum value  $v_M \in V : v_M \geq v_i \quad \forall v_i \in V$ , and be the bit representation of its *pid*  $b_{x-1}b_{x-2} \dots b_2b_1b_0$ . After each time step  $dim$  the processors with *pids* of the form shown in Equation 3.1 have the maximum value because they have received it in any of the previous iterations. After  $x = \log p$  steps all the processors will be within this set.

$$b_{x-1}b_{x-2} \dots b_{dim+2}b_{dim+1} \underbrace{X \dots XX}_{dim + 1 \text{ times}} . \quad (3.1)$$

We can also prove that given any processor  $p$  there exists a chain of *exchanges* less than  $\log p$  that goes from  $p_M$  to  $p$ .

### 3.4 Running Time

Finally, to calculate the running time we have to take into account that both the *exchange* statement and the *if* statement are  $\Theta(1)$  because we are sending before receiving. Since there are  $\log p$  steps the total running time will be  $\Theta(\log p)$  which is the best we can do since there is a logarithmic lower bound as we saw due to the communication constraints.